

Predicting Urban Taxi Demand and Smart Charging Strategies for Electric Fleets

Advanced Analytics and Applications

Team 1



Authors:

Ivan Kamal Mehieddin (7396805)

Nils Bringewatt (7443820)

Maike Katterbach (7444323)

Vitus Groß (7380537)

Contact: mkatter1@smail.uni-koeln.de

Supervisor: Univ.-Prof. Dr. Wolfgang Ketter

Co-Supervisor: Janik Muires

Chair of Information Systems for Sustainable Society
Faculty of Management, Economics and Social Sciences
University of Cologne

2026-04-01

Table of contents

1	Introduction	1
1.1	Problem Description	1
1.2	Business Goal	1
1.3	Data Mining Goal	1
2	Data Description	1
2.1	Taxi Trip Data	2
2.2	Weather Data	2
2.3	Spatial Reference Data	2
2.4	Electricity Price Data	2
2.5	Granularity and Limitations	2
3	Data Preparation	3
4	Descriptive Analytics	3
4.1	Temporal Demand Patterns	4
4.2	Trip Characteristics and Idle Time	4
4.3	Spatial Demand and Resolution	4
4.4	Spatial Hot-Spot Detection	5
5	Predictive Analytics	5
5.1	Problem Description	5
5.2	Modeling Approach	5
5.3	Validation Strategy	6
5.4	Results	6
5.4.1	Support Vector Regression – Community Area, 4-Hour	7
5.4.2	Neural Network – Community Area, 4-Hour	7
5.4.3	Spatial and Temporal Resolution Variations	7
5.5	Discussion	8
5.5.1	Model Comparison	8
5.5.2	Practical Implications	8
6	Intelligent Charging with Reinforcement Learning	8
6.1	Problem Setup	8
6.2	Markov Decision Process	9
6.3	Cost Environments and Electricity Prices	9
6.4	Learning Algorithm and Benchmarks	10
6.5	Results	11
6.6	Operational Implications and Limitations	11
7	Conclusion	12
7.1	Key Findings for Fleet Operations	12
7.2	Demand Forecasting Strategy	12
7.3	Charging Strategy and Reinforcement Learning	12
7.4	Private versus Public Charging Infrastructure	12
7.5	Further Analysis and External Data	13
7.6	Advantages, Limitations, and Business Recommendation	13
8	Appendix: Detailed Tables and Figures	13
8.1	Cleaned Dataset: Column Reference	13
8.2	Descriptive Analytics: Spatial Hot Spots	16

8.3	Predictive Analytics: Metric Definitions	16
8.4	Predictive Analytics: Baseline	17
8.5	Predictive Analytics: Community Area, 4-Hour Resolution	18
8.6	Predictive Analytics: Community Area, 1-Hour Resolution	20
8.7	Predictive Analytics: Census Tract, 4-Hour Resolution	23
9	Citations and References	24

List of Figures

1		16
2	Raw per-window demand is heavily right-skewed, which is why both SVR and NN regress on <code>log1p(count)</code> rather than raw counts.	17
3	Profile baseline predictions vs. actual demand across an evaluation week, highest-demand community area (community area, 4-hour resolution). The baseline tracks the daily and weekly cycle closely.	17
4	Profile baseline test R^2 per community area, 4-hour resolution. The pooled R^2 (0.933) is dominated by the highest-volume areas; the per-area distribution is much wider, with 21 of 77 areas negative.	18
5	Actual, baseline and SVR-variant demand across the March evaluation week, Community Area 8, community area 4-hour resolution.	19
6	Pooled LinearSVR per-area skill vs. the profile baseline, all 77 community areas, 4-hour resolution.	19
7	Actual, baseline and neural-network demand across the March evaluation week, Community Area 8 (highest demand), community area 4-hour resolution.	20
8	Tuned deep NN per-area skill vs. the profile baseline, all 77 community areas, 4-hour resolution. O'Hare (upper left) is the one clearly negative area.	20
9	SVR March-week demand, Community Area 8, 1-hour resolution.	21
10	Neural-network March-week demand, Community Area 8, 1-hour resolution.	21
11	Pooled LinearSVR per-area skill vs. the profile baseline, all 77 community areas, 1-hour resolution.	22
12	Tuned deep NN per-area skill vs. the profile baseline, all 77 community areas, 1-hour resolution.	22
13	Scope map: the 179 of 878 census tracts with enough tract-precision data to support a tract-level model, across 11 community areas.	23
14	Profile baseline test R^2 per census tract, 4-hour resolution. Wider spread than the community-area version (Figure 4), consistent with more units competing for the same total demand.	23
15	Actual, baseline and neural-network demand across the March evaluation week, top-demand census tract (17031760900), 4-hour resolution (estimated trips after inverse-retention scaling).	24
16	Tuned deep NN per-tract skill vs. the profile baseline, 179 in-scope census tracts, 4-hour resolution.	24

1 Introduction

1.1 Problem Description

Urban taxi demand is highly uneven across both space and time. In a city such as Chicago, trip volumes vary by neighborhood, hour of day, weekday, weather conditions, airport activity, commuting patterns, and local points of interest. For a fleet operator, this creates a practical planning problem: vehicles should be available where passengers are likely to request rides, while unnecessary idle time, empty repositioning, and inefficient charging should be reduced.

This project uses Chicago taxi trips as a proxy for urban ride-hailing demand. The core problem is to understand and predict how many trips will start in different parts of the city during a given time window. In this report, demand is therefore defined as the number of taxi pickups per spatial unit and time interval.

1.2 Business Goal

The business goal is to support better operational decisions for a taxi or ride-hailing fleet operator. More accurate knowledge of where and when demand occurs can help operators position vehicles in high-demand areas, reduce passenger waiting times, avoid unnecessary idle time, and improve the utilization of available vehicles. For an electric fleet, this planning task is connected to charging decisions: vehicles need enough energy to serve expected demand, but charging should be scheduled in a cost-efficient way.

From a managerial perspective, the project aims to turn historical trip data into practical recommendations for fleet positioning, demand planning, and charging strategy. The final value of the analysis is not only whether a model predicts demand accurately, but whether the results can inform decisions that improve service reliability and operational efficiency.

1.3 Data Mining Goal

The data mining goal is to build a spatio-temporal demand prediction task from raw taxi trip records and external contextual data. Individual trips are aggregated into panels by time window and spatial unit, and the target variable is the number of pickups in each panel cell. Calendar variables, weather data, and spatial identifiers are used as explanatory features.

The project evaluates whether machine learning models can predict aggregate taxi demand at useful resolutions. In particular, Support Vector Machine models and neural networks are trained and compared against simple time-profile baselines. Model performance is assessed using appropriate error metrics and benchmarking, with attention to whether the models perform well not only overall but also in the high-demand areas that matter most for fleet operations. In a second analytical task, a reinforcement learning approach is used to study how charging decisions for an electric taxi can be optimized under operational constraints.

2 Data Description

The analysis combines taxi trip records, hourly weather observations, spatial reference data, and electricity prices. Taxi and weather data cover 2024-01-01 through 2026-05-31. The period ends in May because the city portal reports recent trips with a delay; including the incomplete June data would create an artificial decline in demand.

2.1 Taxi Trip Data

The main source is the Chicago Taxi Trips dataset from the City of Chicago Data Portal (dataset `Taxi Trips (2024-)`). Each row represents one reported trip by a licensed Chicago taxi. Before cleaning, the extract contains 16,086,196 rows and 21 columns. It includes trip timestamps and duration, distance, fare components, payment type, taxi company, anonymized identifiers, and pickup and dropoff locations.

Locations are reported through the identifiers and centroids of Chicago's 77 community areas and 878 census tracts. Census tracts provide finer spatial detail but are frequently masked in low-volume areas for privacy. Where tract information is unavailable, the coordinate columns may instead contain community-area centroids. This mixed spatial resolution limits fine-grained interpretation and is considered during data preparation and spatial analysis.

2.2 Weather Data

Hourly weather observations are obtained from the Open-Meteo archive API. The dataset contains 21,168 observations and nine variables for the same period as the taxi data. Variables include temperature, apparent temperature, precipitation, snowfall, wind speed, cloud cover, and WMO weather conditions. Weather is joined to the demand data by timestamp and provides exogenous features that can also be known for future prediction periods.

2.3 Spatial Reference Data

Official boundary files provide polygons for Chicago's 77 community areas, 878 census tracts, and the city boundary. They support mapping, spatial joins, and demand aggregation at different geographic resolutions. Point-of-interest data from OpenStreetMap add locations such as transport facilities, amenities, and tourism-related places, helping to characterize areas that may generate taxi demand.

2.4 Electricity Price Data

The reinforcement-learning analysis uses EPEX SPOT day-ahead prices obtained through the Fraunhofer ISE Energy-Charts API. The data cover the German–Luxembourg bidding zone from January through March 2025 and are converted from EUR/MWh to EUR/kWh. For each of the 90 days, prices between 14:00 and 16:00 are transformed into eight 15-minute values matching the charging decisions. These published prices are observable to the price-aware charging agent; the price ultimately billed is simulated by adding a random short-term deviation. Details of this transformation and the simulated price process are given in the reinforcement-learning section.

2.5 Granularity and Limitations

The cleaned taxi data contain 14,465,662 rows and 27 columns. For spatial analytics and predictions, taxi trip observations are aggregated from individual trips into spatio-temporal demand panels, using community areas or census tracts and hourly or four-hour windows. Key limitations are privacy-related location masking, the reporting lag in recent portal data, mixed spatial precision, and the use of licensed taxi trips as a proxy for wider ride-hailing demand.

3 Data Preparation

The raw taxi export contains 16,086,196 rows across 21 columns. The cleaning pipeline concatenates the 2024, 2025, and 2026 CSV exports, applies the steps below, and writes the result out as a single Parquet file. In total, 10.1% of trips (1,620,534 rows) are removed:

1. **Duplicate removal.** Trips are deduplicated on `trip_id` before any other filtering (775 rows removed), keeping the first occurrence. Rows without a `trip_id` are set aside first and rejoined afterward, since `drop_duplicates` would otherwise treat several genuine but ID-less trips as duplicates of each other.
2. **Handling missing spatial identifiers.** Chicago reports location at two resolutions: `census_tract` (fine-grained) and `community_area` (77 coarse zones covering the whole city). Tracts are missing for over half of trips — not randomly, but because the city masks any tract with too few trips, for privacy. Community areas are far more complete (98.2% pickup, 91.6% dropoff), since they're only missing when a trip starts or ends outside city limits. Rather than dropping those out-of-city trips, we keep every row and add `pickup_flag/dropoff_flag` columns marking whether an area is present, so later spatial analyses can filter on the flags without losing the trips for everything else.
3. **Resolution-consistent centroids.** The centroid latitude/longitude columns silently mix resolutions: where a tract exists, they hold its centroid; where it doesn't, they fall back to the community area's centroid. We derive a clean area-to-centroid mapping from exactly the rows where the raw coordinate already is the area centroid and join it back as separate `*_ca_centroid_lat/lon` columns, giving a resolution-consistent coordinate for any area-level aggregation.
4. **Removal of incomplete trip records.** Trips missing `fare`, `trip_total`, `trip_seconds`, or `trip_miles` are dropped (34,687 rows), since no meaningful analysis is possible without them.
5. **Outlier and plausibility filtering.** Negative `tips/tolls/extras` would be removed, though none occur in practice — zero is legitimate, for example for cash trips that report no tip. Trips below Chicago's legal minimum fare (\$3.25 flag-drop) or shorter than 60 seconds are treated as test or aborted entries and removed, together with trips above the 99.9th percentile of distance (43.0 miles) or duration (165 minutes). These bounds account for the bulk of all removed rows (1,585,072 of 1,620,534).

`payment_type` and `company` are left untouched throughout; their missing values remain as NA and are handled per analysis.

The cleaned dataset has 14,465,662 rows and 27 columns, spanning 2024-01-01 to 2026-05-31. Table 1 in the appendix gives the full column-by-column reference: dtype, completeness, distinct-value count, and description for every column in the cleaned file.

4 Descriptive Analytics

The descriptive analysis examines when and where taxi demand occurs, how trips differ in length, duration, and price, and whether a small number of spatial hot spots account for a large share of observed demand. Unless stated otherwise, the results in this section refer to the cleaned dataset of 14,465,662 trips between January 2024 and May 2026.

4.1 Temporal Demand Patterns

Demand follows a pronounced daily cycle. The lowest volume occurs at 03:00, after which the number of pickups increases rapidly during the morning. Demand remains high from late morning through the evening and reaches its maximum at 17:00, consistent with commuting and after-work travel. It then decreases steadily during the evening. Aggregating the data into four-hour intervals smooths short-lived fluctuations while preserving this overall cycle, which supports the use of both hourly and four-hour windows in the predictive analysis.

The weekday profile is equally distinct. Normalized by the number of observed calendar days, Thursday has the highest average demand at approximately 18,856 trips per day, followed by Wednesday at 18,358. Demand is substantially lower on Saturday (13,314) and reaches its weekly minimum on Sunday (12,591). The initially plausible expectation of a strong late-night weekend peak is not supported by the hourly weekend profile. Taxi demand in this dataset is therefore more closely associated with regular weekday mobility, business activity, and commuting than with weekend nightlife alone. Monthly volumes vary as well, but comparisons across years must be interpreted cautiously because 2026 is only observed through May and the underlying portal data have not accumulated uniformly across all months.

4.2 Trip Characteristics and Idle Time

Trip characteristics are strongly right-skewed: most rides are relatively short, while a smaller number of long airport and cross-city journeys raise the averages. The median trip is 3.85 miles long and takes 16.00 minutes, compared with means of 7.09 miles and 20.97 minutes. The median metered fare is \$15.25 and the median total charge is \$17.75; the corresponding means are \$22.06 and \$27.01. This difference between medians and means confirms that a small upper tail of long and expensive trips has a material effect on aggregate revenue and vehicle utilization. For a fleet operator, the median describes a typical assignment more reliably, whereas the mean is more relevant for total energy and capacity planning.

Idle time can in principle be approximated as the interval between a taxi's recorded trip end and its next recorded trip start. The current exploratory notebook does not calculate this measure, and no idle-time estimate is therefore reported here. A valid estimate requires sorting the complete dataset by taxi identifier, calculating consecutive gaps before sampling, and separating genuine waiting time from off-shift periods and incomplete reporting.

4.3 Spatial Demand and Resolution

Pickup demand is substantially more concentrated than a city-wide view would suggest. Among trips with a reported pickup community area, Near North Side accounts for 22.99%, O'Hare for 21.02%, the Loop for 16.72%, and Near West Side for 10.33%. Together, these four areas generate more than 70% of spatially assigned pickups. The pattern links demand to three types of urban activity: the central business and tourism districts, major transport hubs, and airport traffic. Pickup and drop-off density maps show broadly similar central concentrations, while airport trips create distinct peripheral peaks.

The choice of spatial and temporal resolution changes what can be learned from these patterns. Community areas provide stable, interpretable operating zones but conceal substantial within-area variation. Census tracts and H3 hexagons reveal more localized structure, yet the counts become sparse and more sensitive to privacy masking and centroid assignment. Likewise, hourly bins retain short demand peaks but are noisier, whereas four-hour bins provide a more stable signal for planning and prediction. The analysis therefore uses coarse units for robust

city-wide comparison and finer units where local positioning decisions justify the additional variance.

4.4 Spatial Hot-Spot Detection

To identify hot spots more systematically, a Gaussian Mixture Model (GMM) is fitted to the tract-precision pickup locations in the 2025 data. The GMM acts as a parametric spatial kernel density estimator: each component has a center, covariance ellipse, and estimated share of total pickup demand. To avoid interpreting privacy-induced community-area centroids as precise locations, the model uses 2,787,934 tract-resolved pickups at 421 locations inside Chicago. Model comparison indicates an interpretable elbow at six components; a 100-meter covariance floor prevents artificial point masses below the approximately 600-meter empirical spatial resolution of the centroid data. The fitted mixture is shown in Figure 1 (appendix).

The fitted surface confirms that demand is highly concentrated. The five largest components cover approximately 97% of tract-resolved pickups. Three components represent the extended downtown area and jointly account for about 73%: River North/Streeterville (approximately 36%), the Union/Ogilvie station district (24%), and the Loop corridor (14%). O’Hare forms a separate component with roughly 21% of demand, while Midway contributes about 3%. The spatial centers coincide with employment, rail infrastructure, hotels, tourism, nightlife, and airport terminals rather than emerging as arbitrary statistical clusters.

The principal hot-spot locations remain broadly stable between day and night, but their relative importance changes. During the day, the central-business components jointly account for approximately 52% of demand and O’Hare for 18%. At night, O’Hare’s share rises to about 36%, River North retains approximately 26%, and the Loop’s relative share falls. Absolute night-time demand is nevertheless lower. Operationally, the result suggests rebalancing a fleet between a stable set of major zones rather than attempting uniform city-wide coverage: vehicles should be concentrated more strongly around the CBD during the day and shifted toward airports and entertainment districts later in the day.

These findings describe observed and served taxi demand, not all latent ride-hailing demand. Centroid coordinates prevent street-level interpretation, privacy masking understates some low-volume tracts, and taxi regulations may make airports more prominent than they would be for app-based ride-hailing services. The hot spots are therefore appropriate for strategic staging and charging-infrastructure decisions, but not for precise curb-level placement.

5 Predictive Analytics

5.1 Problem Description

We predict taxi pickup demand – the number of trips originating in a spatial unit within a fixed time window. Two model families are used: Support Vector Regression (SVR) and feedforward neural networks (NN). Both are evaluated at community-area resolution in 4-hour and 1-hour windows; census-tract resolution (4-hour) is evaluated with the neural network only. Community area at 4-hour resolution is treated as the default setting throughout, and the other resolutions are presented as deviations from it.

5.2 Modeling Approach

Both families share the same task framing: a single **pooled model per experiment**, trained across all spatial units at once, with spatial identity encoded as one-hot fixed-effect

dummies rather than fit as separate per-unit models. Two families were tried for complementary reasons: SVR (linear and kernelised) is cheap to fit and, in its linear form, interpretable through its coefficients; a feedforward network can in principle learn non-linear interactions between spatial identity and time features that a linear model cannot without hand-built interaction terms.

The central obstacle is that demand is not one distribution but dozens. Per-unit demand varies roughly three orders of magnitude across community areas, and each unit has its own daily and weekly rhythm: an airport runs around the clock on flight schedules, while a downtown office district peaks at commute hours and goes quiet overnight. A pooled model must represent all of that with one set of parameters. This is hardest for a linear SVR, whose single global coefficient vector can only find one compromise direction across every unit at once; a kernel SVR could in principle represent the heterogeneity, but its training cost scales quadratically to cubically in the number of rows, which forces the kernel variants onto a 10,000-row subsample.

Both use the same exogenous feature set: weather (temperature, precipitation, snowfall, wind, cloud cover), cyclical time encodings (sine/cosine pairs for hour-of-day and weekday, two harmonics each), month, and weekend/holiday indicators. Per the constraints of this course, no lagged or autoregressive demand features are used. Both regress on $\log_{1p}(\text{count})$, with predictions back-transformed via a clipped $\exp m1$, since raw demand is heavily right-skewed (Figure 2).

5.3 Validation Strategy

Every model is judged against a **profile baseline**: the train-set mean demand for each (hour-of-day, weekday, month) bucket, computed separately for each spatial unit. Because it is fitted per unit, the baseline sidesteps the heterogeneity problem described above – it never has to compromise one parameter set across areas whose demand differs by three orders of magnitude – which is what makes it a demanding reference throughout this section. Both model families use the same chronological holdout: training on 2024-01-01 through 2025-12-31 and testing on 2026-01-01 through 2026-05-31, over a complete zero-filled grid of every unit crossed with every time window, so that evaluation includes genuine zero-demand periods and not only observed trips.

The baseline already explains a large share of the variation in the data. At the default community-area, 4-hour resolution its pooled test R^2 is 0.933 (Table 2), and it tracks the daily and weekly demand cycle closely across an evaluation week (Figure 3); the same holds at the finer settings (pooled R^2 0.918 at 1-hour and 0.896 at census-tract level). Per-unit fit is far more variable – median per-area R^2 is only 0.208 at 4-hour resolution, with 21 of 77 areas negative (Figure 4) – because the pooled figure is dominated by the few highest-volume units.

Two headline metrics summarise each model; both are defined formally in Section 8. **Demand-weighted skill** is a per-unit ratio of the model’s MAE to the baseline’s MAE (Equation 1), aggregated with weights proportional to each unit’s median demand (Equation 2) so that 0 = tied with the baseline, positive = better, negative = worse; it cancels each unit’s own scale and is comparable across resolutions. **Weighted RMSE** (Equation 3) measures error magnitude in trips per window, with high-volume rows up-weighted, and is the metric used to compare the two model families directly.

5.4 Results

Community area at 4-hour resolution is the default setting; the other settings follow as deviations from it.

5.4.1 Support Vector Regression – Community Area, 4-Hour

Ten SVR configurations were run at this resolution (Table 3). A pooled **LinearSVR** reaches demand-weighted skill -0.827 (weighted RMSE 254.94); the sample-weighted and demand-tiered variants stay strongly negative. Kernel SVR tuned on a 10,000-row subsample scores higher (tuned RBF -0.232 , polynomial -0.221) but remains below the baseline. Fully separate per-area models, fit only for the 20 areas above a median-demand threshold of 10 trips/window, score higher still: the tuned per-area RBF SVR is the single configuration with positive demand-weighted skill, $+0.012$ (weighted RMSE 143.80). Figure 5 shows the pooled variants over the March week in Community Area 8, and Figure 6 maps pooled **LinearSVR** skill across the 77 areas. All ten configurations are collected in Table 3.

5.4.2 Neural Network – Community Area, 4-Hour

Two networks were run: an untuned shallow network (one hidden layer) and a deeper network selected by a one-factor-at-a-time architecture search. Their test scores are in Table 4. The shallow network reaches demand-weighted skill -0.028 (weighted RMSE 144.96) and the tuned deep network -0.027 (weighted RMSE 145.93), against the baseline’s 139.02; both sit just below the baseline on skill. Figure 7 shows the three models tracing demand across the March evaluation week in the highest-demand area (Community Area 8). The per-area skill map (Figure 8) is at or above zero across most of the city, with O’Hare (community area 76) the one clearly negative area (skill -0.227).

5.4.3 Spatial and Temporal Resolution Variations

Community area, 1-hour. At 1-hour resolution the same 77 areas are split across 24 hourly buckets instead of 6, quadrupling the rows and the baseline’s granularity. Test scores for both families are in Table 5. For SVR, demand-weighted skill of the pooled **LinearSVR** falls to -1.016 and the best tuned kernel (RBF) reaches -0.331 . For the neural network, the tuned deep configuration selected at 4-hour resolution (reused unchanged) reaches demand-weighted skill $+0.062$ (weighted RMSE 36.26, against the baseline’s 39.46) – the only positive NN result in the study – while the shallow network stays negative (-0.048). March-week traces appear in Figure 9 and Figure 10, and per-area skill maps in Figure 11 and Figure 12; scores for both families are in Table 5.

Census tract, 4-hour. Only 44.6% of trips carry a true tract-level location; Chicago masks the rest to the community-area centroid whenever a tract had too few trips in a period, for privacy. To recover realistic tract-level demand, tract-precision coverage was first measured per community area, and only areas retaining at least 5% of their trips at tract precision were kept – a natural gap in the data separates areas at 4.85% from those at 7.17% – leaving **179 of 878 tracts across 11 community areas** in scope (Figure 13); the remaining areas fall back to the community-area model. Within scope, per-window counts were built from tract-precision trips only and then scaled up by each area’s inverse retention rate, so the observed fraction is extrapolated to an estimate of each tract’s total demand. This keeps the masked points from spiking demand artificially onto whichever tract holds a community area’s centroid.

Only neural networks were run at this resolution (Table 6). The shallow network reaches demand-weighted skill -0.098 (weighted RMSE 113.77, against the baseline’s 100.99) and the deep network -0.340 . March-week traces for the top tract are in Figure 15 and the per-tract skill map in Figure 16. SVR was not run at tract level because a per-tract one-hot encoding (up to 878 dummy variables) is computationally infeasible for kernel SVR.

5.5 Discussion

5.5.1 Model Comparison

Across every setting the profile baseline is a strong reference, and it is hardest to beat exactly where demand is highest. Both families lose demand-weighted skill specifically in the busiest areas, because a single pooled parameter set cannot give a high-volume downtown its own scale and rush-hour shape while also serving the low-demand majority, whereas the per-unit baseline never faces that compromise. The neural network has the representational capacity to learn per-area time interactions and comes closest: near break-even at 4-hour resolution and clearly positive (+0.062) once the same architecture is reused at 1-hour resolution. Pooled SVR trails further behind at every resolution, and only abandoning pooling altogether – a dedicated per-area RBF model for the 20 busiest areas – yields a positive result (+0.012). The two families’ failures also share a location: the deep NN’s worst area at 4-hour resolution is O’Hare, whose round-the-clock, flight-driven demand differs in *shape*, not just scale, from the commute cycle the pooled model averages toward.

Resolution matters in opposite directions for the two families. For matched SVR types, skill is worse at 1-hour than at 4-hour, so the finer grid sharpens the high-demand gap; for the neural network, the 4-hour architecture transfers to 1-hour cleanly and produces its one clear win. Census-tract results are the weakest overall: even with the inverse-retention extrapolation, the per-unit signal is thin once 179 units compete for demand, and both networks fall below the tract baseline.

5.5.2 Practical Implications

The consistent pattern – pooled models fail in the high-demand areas that matter most for allocation – points toward a targeted deployment rather than one universal model. A fleet operator does not need accurate predictions for every low-demand area; those are served about as well by the cheap profile baseline as by anything tested here. The value concentrates in the small set of highest-demand areas, and the only configurations that beat the baseline there – the per-area RBF SVR and the 1-hour deep NN – are precisely the ones that either abandon pooling or exploit a finer temporal grid. The practical recommendation is therefore to fit dedicated models for the handful of high-value areas and leave the long tail of low-demand areas to the profile baseline, which already handles them about as well as any model tried here.

6 Intelligent Charging with Reinforcement Learning

6.1 Problem Setup

The task is defined around an electric taxi that arrives at its home charger at 14:00 and departs at 16:00. During this two-hour window, a charging decision is made every 15 minutes, resulting in eight decisions per episode. The objective is to minimize charging cost while leaving enough energy for the next shift. Any energy shortage at departure is penalized at EUR 100 per missing kWh – an arbitrarily set severity, chosen to make a shortfall categorically worse than any charging cost rather than derived from an actual operational figure.

To represent a heterogeneous fleet, battery capacity is sampled uniformly from 40, 60, and 77 kWh in each episode, corresponding to compact, mid-size, and large vehicle classes. Initial state of charge (SOC) is independently sampled between 20% and 60% of capacity. The energy required for the next shift is drawn from $d \sim \mathcal{N}(30, 5^2)$ kWh and clipped to the vehicle’s capacity; this value enters the shortage penalty only at the terminal step rather than being observed during

charging. The demand distribution is therefore known throughout charging, while the specific requirement remains unresolved until the terminal step and must be hedged against.

6.2 Markov Decision Process

The charging problem is modeled as a finite-horizon Markov Decision Process. Without day-ahead price information, the state is

$$s_t = [SOC_t, t, C],$$

where SOC_t is the current energy level, $t \in \{0, \dots, 8\}$ is the time step, and C is battery capacity. The realized demand is deliberately excluded from the state. When day-ahead price information is included, the state is extended with the current day-ahead price and the mean remaining day-ahead price:

$$s_t^{PEX} = [SOC_t, t, C, \pi_t^{DA}, \bar{\pi}_{\geq t}^{DA}].$$

Four discrete charging powers of 0, 5, 11, and 22 kW are used initially. Because this coarse spacing leads to degenerate behavior under the exponential cost function, a refined grid of 0, 4, 8, 12, 16, 20, and 22 kW is used instead in a second experiment. Both DQNs are retrained from scratch on this grid, which is used for the primary results below. For a 15-minute interval, the transition is

$$SOC_{t+1} = \min(C, SOC_t + p_t \cdot 0.25).$$

The step reward is the negative charging cost. After the eighth decision, a terminal penalty is added if $SOC_8 < d$:

$$r_8 = -\text{cost}_8 - 100 \max(0, d - SOC_8).$$

A cost–reliability trade-off results: additional energy is inexpensive relative to a shortage, while the tail risk must be learned from sampled experience.

6.3 Cost Environments and Electricity Prices

Two **gymnasium** environments are used, sharing the same battery dynamics and hidden-demand process but differing in charging-cost model. In the stylized environment, the cost of selecting power p in interval t is defined as

$$\text{cost}(t, p) = \alpha_t \exp(p),$$

where α_t increases from 2.0×10^{-10} to 7.2×10^{-10} over the charging window. The small coefficients are required because $\exp(22)$ is approximately 3.6×10^9 . Even charging at 22 kW in all eight intervals costs only 11.90 reward units, compared with the shortage penalty of EUR 100 per missing kWh. Consequently, the original costs of 0, 5, and 11 kW are almost indistinguishable, while 22 kW constitutes a separate, much more expensive tier. Jumps between adjacent actions are reduced by the refined grid, allowing more graduated control, although the cost function remains nonlinear. The stylized model should therefore be interpreted as an artificial power-dependent reward rather than a realistic electricity tariff.

Ninety daily price profiles from the German–Luxembourg bidding zone for January through March 2025 are used in the EPEX environment. Hourly EPEX SPOT day-ahead prices are converted from EUR/MWh to EUR/kWh and linearly interpolated to the eight quarter-hour decisions between 14:00 and 16:00. The resulting day-ahead prices range from EUR -0.026 to EUR 0.371 per kWh, with a median of EUR 0.110 per kWh.

Only the day-ahead price – the current value and the mean of the remaining window, published one day in advance – is observed by the price-aware agent. Charging is billed at the realized price

$$\pi_t = \pi_t^{DA} + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, 0.015^2) \text{ EUR/kWh},$$

which is not observable before the decision. ε_t represents the gap between the day-ahead forecast and the delivery price and is simulated rather than drawn from a second price dataset. Price profiles are loaded from a local cache, with the Fraunhofer ISE Energy-Charts API used only as a fallback.

Both cost formulations deviate from a purely quadratic p^2 charging-cost specification defined above: the stylized cost emphasizes a strongly nonlinear power penalty, while the EPEX cost treats energy procurement as the relevant marginal cost. The two environments are therefore complementary experiments rather than directly comparable tariffs, and a policy learned under the stylized cost structure cannot be expected to remain cost-efficient under real per-kWh prices.

6.4 Learning Algorithm and Benchmarks

Two Deep Q-Networks (DQN) are trained: a price-unaware agent in the stylized environment and a price-aware agent in the EPEX environment. A multilayer perceptron with two hidden layers of 128 units and identical hyperparameters is used for both. Training lasts 200,000 time steps, equivalent to approximately 25,000 simulated charging days. Performance is evaluated on 1,000 paired episodes, using identical seeds so that capacity, initial SOC, demand, price day, and simulated price deviations match across strategies.

The operational benchmark is a price-blind greedy heuristic. Because demand is unknown, charging is not targeted at the realized requirement; instead, a fixed target is used:

$$\tau = \min(C, \mu + 2\sigma) = \min(C, 40 \text{ kWh}),$$

corresponding to an approximate 97.7% service level under the assumed normal distribution. At each step, the energy still missing from this target is spread over the remaining time, and the smallest sufficient charging power is selected.

An exact brute-force benchmark is used, evaluating every action sequence. For the original grid this requires $4^8 = 65,536$ sequences, evaluated on 100 episodes. The refined grid increases the search to $7^8 = 5,764,801$ sequences, so the EPEX oracle evaluation is restricted to 40 episodes. The realized demand and, in the EPEX setting, the realized billed prices, are known to this oracle benchmark in advance, making it a theoretical upper bound rather than an implementable policy.

6.5 Results

An action-space problem is evident in the original four-action experiment. In the stylized environment, the price-unaware DQN succeeds in 95.5% of the 1,000 episodes and achieves a mean reward of -17.201, compared with 98.9% and -2.581 for greedy. Because charging at 11 kW is almost free relative to 22 kW, 11 kW is selected almost exclusively by the learned policy; its action entropy is only 0.039 bits.

With the refined grid, the price-unaware DQN succeeds in 98.3% of episodes and achieves a mean reward of -1.855. Greedy succeeds in 98.9% and obtains -2.058. The DQN’s action entropy rises to 1.522 bits, indicating graduated control rather than collapse onto a single action. Much of the original performance gap is therefore explained by action-grid resolution: in its training environment, the reliability of greedy is nearly matched by the refined DQN, whose mean reward is slightly higher.

Under EPEX prices, greedy succeeds in 98.9% of episodes, compared with 98.3% for the refined DQN trained in the stylized environment and 75.2% for the refined price-aware DQN. Their unconditional mean charging costs are EUR 1.7223, EUR 2.0424, and EUR 0.9931 per day, respectively. Mean rewards, which also include shortage penalties, are -3.7775 for greedy, -3.8872 for the DQN without EPEX data, and -88.0377 for the price-aware DQN. The price-unaware DQN now nearly matches greedy in reliability, but its charging cost remains 18.8% higher on episodes in which both succeed, because its policy was trained against the stylized exponential cost rather than real per-kWh prices.

In the 752 episodes where both greedy and the price-aware DQN cover demand, the price-aware DQN saves EUR 0.6296 per day, or 41.7%, relative to greedy and is cheaper in 88% of episodes. Policy maps show the intended behavior: at sufficiently high SOC, high power is selected during cheap intervals and charging is deferred during expensive ones; as SOC decreases or departure approaches, charging urgency increasingly dominates over price. However, this conditional comparison does not establish that the price-aware DQN is the better overall policy: even among successful episodes, different quantities of reserve energy may be purchased by each strategy, and the much lower overall success rate remains excluded from this cost figure.

The cost advantage remains inseparable from inadequate demand insurance: a 75.2% success rate is not operationally acceptable for a taxi fleet, even when costs are lower on successful days. On the reduced 40-episode sample, the oracle reaches 100% success at approximately EUR 0.87 per day, but this relies on information unavailable to any real agent. Day-ahead information therefore has value for charge timing, while the price-aware DQN still requires an explicit reliability safeguard.

6.6 Operational Implications and Limitations

In the stylized experiment, the refined DQN and the service-level heuristic perform similarly, while the sensitivity of DQN to action-space design is demonstrated by the coarse-grid result. Under dynamic prices, cheap intervals can be exploited by learned control, particularly when the vehicle arrives with sufficient SOC and prices vary substantially within the charging window. However, part of the apparent saving of the price-aware DQN is attributable to carrying too little reserve. The learned policy should therefore be placed behind a safety layer in a practical implementation, for example a hard minimum-SOC rule during the final intervals, or trained with a stronger asymmetric shortage objective. Private charging combined with a dynamic electricity contract is the most controllable setting for realizing timing-based savings; less opportunity for price-aware scheduling exists under public charging with fixed tariffs.

The analysis remains a proof of concept. One two-hour charging window, stylized battery dynamics, no charging losses or battery degradation, and a fixed normal demand distribution are considered. A fixed-power flat-rate baseline has not yet been evaluated, and the required sensitivity analysis for alternative demand levels such as 25, 30, and 35 kWh is still missing. Moreover, the linear energy-cost model should either be complemented by the assignment’s quadratic power-cost formulation or retained only with the stated interpretation. These additions are necessary before an infrastructure or deployment recommendation can be made.

7 Conclusion

7.1 Key Findings for Fleet Operations

The analysis shows that the greatest immediate value for a potential fleet operator lies in combining robust demand patterns with transparent operational rules rather than relying on a single complex model. Taxi demand is strongly concentrated in both space and time: more than 70% of spatially assigned pickups originate in Near North Side, O’Hare, the Loop, and Near West Side, while demand peaks on weekdays and in the late afternoon. These patterns support a zone-based operating strategy in which vehicles are staged primarily in the central business district and at major transport hubs, with time-dependent rebalancing toward airports and entertainment areas later in the day. The results should, however, be interpreted as observed taxi demand rather than total latent mobility demand.

7.2 Demand Forecasting Strategy

For demand forecasting, model complexity did not consistently translate into operationally better predictions. The historical profile baseline remained difficult to outperform, particularly in high-volume areas. Most pooled SVR and neural-network specifications produced negative demand-weighted skill, although the deep neural network achieved a small positive skill at one-hour resolution and dedicated RBF models for the 20 highest-demand community areas also slightly improved on the baseline. This suggests a hybrid forecasting architecture: use simple profile forecasts for the long tail of low-volume areas and train dedicated models only for the small number of strategically important high-demand zones. Such an approach is easier to maintain, more interpretable, and better aligned with the concentration of fleet-allocation value than one city-wide pooled model.

7.3 Charging Strategy and Reinforcement Learning

The charging experiment leads to a similar conclusion. The price-blind greedy policy reached a 98.9% success rate and clearly outperformed both DQN agents once shortage penalties were included. The price-aware DQN demonstrated that dynamic electricity prices can reduce charging costs when sufficient energy is available, but its 75.9% reliability is not acceptable for fleet operations. Reinforcement learning should therefore not yet control charging without a safety mechanism. A practical system could use price-aware scheduling within hard minimum-SOC and departure-energy constraints, preserving the cost benefit of flexible charging without putting the next shift at risk.

7.4 Private versus Public Charging Infrastructure

For a fleet with regular depot returns, private charging infrastructure is the preferable long-term option. It provides predictable availability, operational control, and access to dynamic

electricity tariffs, while public chargers introduce travel time, queues, uncertain availability, and generally less control over prices. Public charging nevertheless remains useful as a backup and for vehicles that cannot reliably return to a depot. The recommended strategy is therefore a mixed system: install private chargers at one or more fleet bases near the main operating corridors and use public fast charging for contingencies and peak demand. The present analysis supports this operational direction, but it does not by itself justify a specific number, capacity, or location of chargers because infrastructure costs, grid-connection limits, depot locations, public-charger availability, and complete vehicle duty cycles are not included in the dataset.

7.5 Further Analysis and External Data

Several extensions would strengthen the business case. The existing trip data could be used to construct full vehicle duty cycles and valid idle-time distributions, estimate repositioning distances, test rolling-origin forecasts, and evaluate prediction performance under a common scope and baseline. Charging simulations should cover complete operating days, alternative energy-demand scenarios, charging losses, battery degradation, charger congestion, and a fixed-power or flat-rate benchmark. Particularly useful external sources would include public charger locations, power levels, prices and occupancy; depot and electricity-grid data; real-time traffic and road-network travel times; flight and public-transport schedules; major events; and more complete weather forecasts. These additions would permit a joint location and capacity analysis comparing the total cost of private infrastructure with public-charging fees, detours, waiting time, and lost service revenue.

7.6 Advantages, Limitations, and Business Recommendation

Overall, the project provides a useful proof of concept rather than a deployment-ready optimization system. Its main advantages are the large trip dataset, the combination of spatial, predictive, and charging analyses, and the use of operationally meaningful benchmarks. Its main limitations are spatial privacy aggregation, incomplete and partly inconsistent model comparisons, the absence of valid idle-time estimates, and simplified charging assumptions. Accordingly, the fleet operator should initially deploy a conservative, interpretable planning system based on demand profiles, hotspot-specific forecasts, and rule-based charging, while collecting the operational and infrastructure data required for a later, safely constrained optimization solution.

8 Appendix: Detailed Tables and Figures

Tables and figures below are referenced from the main text and grouped by chapter, in the same order as they are cited. All Predictive Analytics figures and score tables – the March-week model comparisons, the per-area skill maps, and the per-resolution score tables – are collected here and pointed to from the Results and Validation Strategy sections.

8.1 Cleaned Dataset: Column Reference

Every column of `data/chicago_taxi_2024_2026_clean.parquet` (14,465,650 rows), with dtype, completeness, and distinct-value count as measured on that file, plus what the column means and any caveat carried over from the cleaning steps in Section 3 (e.g. census-tract masking, the two coordinate resolutions).

Table 1: Data dictionary, cleaned dataset.

Column	Dtype	Non-Null %	Distinct	Description
trip_id	str	100.0%	14,465,650	Unique trip identifier (hash). Deduplicated in cleaning step 1.
taxi_id	str	100.0%	3,354	Anonymized vehicle identifier (hash).
trip_start_timestamp	date-time64[us]	100.0%	84,660	Local Chicago date/time the trip started.
trip_end_timestamp	date-time64[us]	100.0%	84,664	Local Chicago date/time the trip ended.
trip_seconds	float32	100.0%	9,722	Trip duration in seconds (≥ 60 s, capped at the 99.9th percentile).
trip_miles	float32	100.0%	4,303	Trip distance in miles (> 0 , capped at the 99.9th percentile).
pickup_census_tract	float32	44.6%	149	Pickup census tract (FIPS code). Missing for ~55% of trips – privacy masking of low-volume tracts.
dropoff_census_tract	float32	43.3%	175	Dropoff census tract (FIPS code). Same privacy masking as pickup.
pickup_community_area	Int64	98.2%	77	One of 77 official community areas for the pickup. Missing means the pickup was outside city limits.
dropoff_community_area	Int64	91.6%	77	Community area for the dropoff. Missing means the dropoff was outside city limits.
fare	float32	100.0%	9,078	Metered fare in USD ($\geq \$3.25$ flag-drop minimum).
tips	float32	100.0%	3,619	Tip amount in USD. 0 is valid – cash trips report no tip.
tolls	float32	100.0%	560	Toll charges in USD. 0 is valid – most trips use no toll roads.
extras	float32	100.0%	3,624	Additional surcharges in USD. 0 is valid.
trip_total	float32	100.0%	14,616	Total amount charged in USD ($\geq \$3.25$).
payment_type	category	100.0%	7	Payment method (Credit Card, Cash, Mobile, Prcard, Unknown, ...). Left untouched by cleaning.
company	str	100.0%	45	Taxi affiliation / operating company. Left untouched by cleaning.
pickup_centroid_latitude	float32	98.2%	571	Pickup latitude: census tract centroid where known, else community area centroid (mixed resolution).
pickup_centroid_longitude	float32	98.2%	565	Pickup longitude, same mixed-resolution rule as latitude.
dropoff_centroid_latitude	float32	92.1%	684	Dropoff latitude, same mixed-resolution rule.
dropoff_centroid_longitude	float32	92.1%	678	Dropoff longitude, same mixed-resolution rule.
pickup_flag	int64	100.0%	2	1 if pickup_community_area is present, 0 if missing (outside Chicago).

Column	Dtype	Non-Null %	Distinct	Description
dropoff_flag	int64	100.0%	2	1 if dropoff_community_area is present, 0 if missing (outside Chicago).
pickup_ca_c entroid_lat	float32	98.2%	77	Pickup community area centroid latitude – always CA-level resolution (added in step 3).
pickup_ca_c entroid_lon	float32	98.2%	77	Pickup community area centroid longitude – always CA-level resolution.
dropoff_ca_c entroid_lat	float32	91.6%	77	Dropoff community area centroid latitude – always CA-level resolution.
dropoff_ca_c entroid_lon	float32	91.6%	77	Dropoff community area centroid longitude – always CA-level resolution.

8.2 Descriptive Analytics: Spatial Hot Spots




Figure 1

8.3 Predictive Analytics: Metric Definitions

The two headline metrics used throughout the Predictive Analytics chapter are defined below. The per-unit **skill score** compares a model's mean absolute error to the profile baseline's for the same spatial unit i ,

$$\text{skill}_i = 1 - \frac{\text{MAE}_{\text{model}, i}}{\text{MAE}_{\text{baseline}, i}} \quad (1)$$

so that 0 means tied with the baseline, 1 means zero error, and a negative value means worse than the baseline. The reported **demand-weighted skill** aggregates the per-unit skills with weights equal to each unit’s median demand m_i , so the highest-volume units dominate,

$$\text{skill}_{\text{demand-weighted}} = \frac{\sum_i m_i \text{skill}_i}{\sum_i m_i}. \quad (2)$$

The **weighted RMSE** is a pooled, absolute-scale error over every test-set row j (a unit crossed with a time window), with y_j and $\hat{y}_j$ the actual and predicted demand and $\bar{d}_{u(j)}$ the train-period mean demand of the unit that row j belongs to,

$$\text{weighted RMSE} = \sqrt{\frac{\sum_j \bar{d}_{u(j)} (y_j - \hat{y}_j)^2}{\sum_j \bar{d}_{u(j)}}}. \quad (3)$$

8.4 Predictive Analytics: Baseline




Figure 2: Raw per-window demand is heavily right-skewed, which is why both SVR and NN regress on $\log_{1p}(\text{count})$ rather than raw counts.




Figure 3: Profile baseline predictions vs. actual demand across an evaluation week, highest-demand community area (community area, 4-hour resolution). The baseline tracks the daily and weekly cycle closely.




Figure 4: Profile baseline test R^2 per community area, 4-hour resolution. The pooled R^2 (0.933) is dominated by the highest-volume areas; the per-area distribution is much wider, with 21 of 77 areas negative.

Table 2: Profile baseline test scores by resolution.

Resolution	Units	Pooled R^2	Weighted RMSE	Median per-unit R^2	Units with $R^2 < 0$
Community area, 4-hour	77	0.933	139.02	0.208	21 of 77
Community area, 1-hour	77	0.918	39.46	0.028	36 of 77
Census tract, 4-hour	179	0.896	100.99	0.195	56 of 179

8.5 Predictive Analytics: Community Area, 4-Hour Resolution

Table 3: SVR test scores, community area, 4-hour resolution. Per-area models cover only the 20 areas above the demand threshold.

Model	Units	Dem.- weighted skill	Median skill	Weighted RMSE	Pooled R^2
Profile baseline	77	0.000	0.000	139.02	0.933
LinearSVR, pooled	77	-0.827	-0.017	254.94	0.778
LinearSVR, demand-weighted	77	-0.675	-0.052	231.91	0.814
LinearSVR, demand-tiered	77	-0.732	0.000	243.60	0.800
RBF, tuned (10k subsample)	77	-0.232	-0.027	176.93	0.892
Polynomial, tuned (10k subsample)	77	-0.221	-0.049	171.71	0.896
LinearSVR, per-area	20	-0.112	-0.091	156.93	0.905
LinearSVR, per-area, tuned	20	-0.125	-0.116	161.10	0.901
RBF, per-area	20	-0.135	-0.105	163.37	0.897
RBF, per-area, tuned	20	+0.012	-0.002	143.80	0.921




Figure 5: Actual, baseline and SVR-variant demand across the March evaluation week, Community Area 8, community area 4-hour resolution.




Figure 6: Pooled LinearSVR per-area skill vs. the profile baseline, all 77 community areas, 4-hour resolution.

Table 4: Neural network test scores, community area, 4-hour resolution.

Model	Dem.- weighted skill	Median skill	Weighted RMSE	Pooled R^2	Fit time
Profile baseline	0.000	0.000	139.02	0.933	1s
Shallow NN (untuned)	-0.028	+0.097	144.96	0.925	324s
Deep NN (tuned)	-0.027	+0.071	145.93	0.927	2,333s




Figure 7: Actual, baseline and neural-network demand across the March evaluation week, Community Area 8 (highest demand), community area 4-hour resolution.




Figure 8: Tuned deep NN per-area skill vs. the profile baseline, all 77 community areas, 4-hour resolution. O'Hare (upper left) is the one clearly negative area.

8.6 Predictive Analytics: Community Area, 1-Hour Resolution

Table 5: Test scores for both families, community area, 1-hour resolution.

Model	Dem.-weighted skill	Median skill	Weighted RMSE	Pooled R^2
Profile baseline	0.000	0.000	39.46	0.918
SVR –	-1.016	+0.022	76.96	0.687
LinearSVR, pooled				
SVR – RBF, tuned	-0.331	+0.015	54.53	0.835

Model	Dem.-weighted skill	Median skill	Weighted RMSE	Pooled R^2
SVR – Polynomial, tuned	-0.361	0.000	56.14	0.793
NN – Shallow (untuned)	-0.048	+0.118	42.28	0.903
NN – Deep (tuned)	+0.062	+0.090	36.26	0.928




Figure 9: SVR March-week demand, Community Area 8, 1-hour resolution.




Figure 10: Neural-network March-week demand, Community Area 8, 1-hour resolution.




Figure 11: Pooled LinearSVR per-area skill vs. the profile baseline, all 77 community areas, 1-hour resolution.




Figure 12: Tuned deep NN per-area skill vs. the profile baseline, all 77 community areas, 1-hour resolution.

Table 6: Neural network test scores, census tract, 4-hour resolution (179 in-scope tracts).

Model	Dem.-weighted skill	Weighted RMSE	Pooled R^2
Profile baseline	0.000	100.99	0.896
Shallow NN (untuned)	-0.098	113.77	0.868
Deep NN (reused 4h config)	-0.340	138.71	0.821




Figure 15: Actual, baseline and neural-network demand across the March evaluation week, top-demand census tract (17031760900), 4-hour resolution (estimated trips after inverse-retention scaling).




Figure 16: Tuned deep NN per-tract skill vs. the profile baseline, 179 in-scope census tracts, 4-hour resolution.

9 Citations and References

You can easily cite your sources using BibLaTeX. For example, you can cite the Quarto documentation (Team, 2024) or a textbook like “R for Data Science” (Wickham et al., 2023). The bibliography will be automatically generated at the end of the document.

References

Team, Q. (2024). Quarto: An open-source scientific and technical publishing system. *Journal of Modern Data Science*. <https://quarto.org>

Wickham, H., Çetinkaya-Rundel, M., & Grolemund, G. (2023). *R for data science*. O'Reilly Media.